

MACRO ECONOMIC DATA

OBJECTIVE:

To understand and analyze the big picture of an economy by studying key indicators and trends, such as GDP growth, inflation, unemployment, and trade, to make informed decisions and policies that support economic well-being and growth.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv(r'/content/macrodata.csv_.csv')
df
```

|     | Date       | year | quarter | realgdp   | realcons | realinv  | realgovt | realdpi |     |
|-----|------------|------|---------|-----------|----------|----------|----------|---------|-----|
| 0   | 3/31/1959  | 1959 | 1       | 2710.349  | 1707.4   | 286.898  | 470.045  | 1886.9  | 2   |
| 1   | 6/30/1959  | 1959 | 2       | 2778.801  | 1733.7   | 310.859  | 481.301  | 1919.7  | 2   |
| 2   | 9/30/1959  | 1959 | 3       | 2775.488  | 1751.8   | 289.226  | 491.260  | 1916.4  | 2   |
| 3   | 12/31/1959 | 1959 | 4       | 2785.204  | 1753.7   | 299.356  | 484.052  | 1931.3  | 2   |
| 4   | 3/31/1960  | 1960 | 1       | 2847.699  | 1770.5   | 331.722  | 462.199  | 1955.5  | 2   |
| ... | ...        | ...  | ...     | ...       | ...      | ...      | ...      | ...     | ... |
| 198 | 9/30/2008  | 2008 | 3       | 13324.600 | 9267.7   | 1990.693 | 991.551  | 9838.3  | 21  |
| 199 | 12/31/2008 | 2008 | 4       | 13141.920 | 9195.3   | 1857.661 | 1007.273 | 9920.4  | 21  |
| 200 | 3/31/2009  | 2009 | 1       | 12925.410 | 9209.2   | 1558.494 | 996.287  | 9926.4  | 21  |
| 201 | 6/30/2009  | 2009 | 2       | 12901.504 | 9189.0   | 1456.678 | 1023.528 | 10077.5 | 21  |
| 202 | 9/30/2009  | 2009 | 3       | 12990.341 | 9256.0   | 1486.398 | 1044.088 | 10040.6 | 21  |

203 rows × 15 columns

```
df.head()
```

|   | Date       | year | quarter | realgdp  | realcons | realinv | realgovt | realdpi | cpi   |
|---|------------|------|---------|----------|----------|---------|----------|---------|-------|
| 0 | 3/31/1959  | 1959 | 1       | 2710.349 | 1707.4   | 286.898 | 470.045  | 1886.9  | 28.98 |
| 1 | 6/30/1959  | 1959 | 2       | 2778.801 | 1733.7   | 310.859 | 481.301  | 1919.7  | 29.15 |
| 2 | 9/30/1959  | 1959 | 3       | 2775.488 | 1751.8   | 289.226 | 491.260  | 1916.4  | 29.35 |
| 3 | 12/31/1959 | 1959 | 4       | 2785.204 | 1753.7   | 299.356 | 484.052  | 1931.3  | 29.37 |
| 4 | 3/31/1960  | 1960 | 1       | 2847.699 | 1770.5   | 331.722 | 462.199  | 1955.5  | 29.54 |

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 203 entries, 0 to 202
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        203 non-null   object
1   year        203 non-null   int64
2   quarter     203 non-null   int64
3   realgdp     203 non-null   float64
4   realcons    203 non-null   float64
5   realinv     203 non-null   float64
6   realgovt    203 non-null   float64
7   realdpi     203 non-null   float64
8   cpi         203 non-null   float64
9   m1          203 non-null   float64
10  tbilrate    203 non-null   float64
11  unemp       203 non-null   float64
12  pop         203 non-null   float64
13  infl        203 non-null   float64
14  realint     203 non-null   float64
dtypes: float64(12), int64(2), object(1)
memory usage: 23.9+ KB
```

```
df['Date'] = pd.to_datetime(df['Date'])
```

```
df=df.set_index(['Date'])
```

df

|            | year | quarter | realgdp   | realcons | realinv  | realgovt | realdpi | cpi     |     |
|------------|------|---------|-----------|----------|----------|----------|---------|---------|-----|
| Date       |      |         |           |          |          |          |         |         |     |
| 1959-03-31 | 1959 | 1       | 2710.349  | 1707.4   | 286.898  | 470.045  | 1886.9  | 28.980  | 13  |
| 1959-06-30 | 1959 | 2       | 2778.801  | 1733.7   | 310.859  | 481.301  | 1919.7  | 29.150  | 14  |
| 1959-09-30 | 1959 | 3       | 2775.488  | 1751.8   | 289.226  | 491.260  | 1916.4  | 29.350  | 14  |
| 1959-12-31 | 1959 | 4       | 2785.204  | 1753.7   | 299.356  | 484.052  | 1931.3  | 29.370  | 14  |
| 1960-03-31 | 1960 | 1       | 2847.699  | 1770.5   | 331.722  | 462.199  | 1955.5  | 29.540  | 13  |
| ...        | ...  | ...     | ...       | ...      | ...      | ...      | ...     | ...     | ... |
| 2008-09-30 | 2008 | 3       | 13324.600 | 9267.7   | 1990.693 | 991.551  | 9838.3  | 216.889 | 147 |
| 2008-12-31 | 2008 | 4       | 13141.920 | 9195.3   | 1857.661 | 1007.273 | 9920.4  | 212.174 | 157 |


```
data = pd.DataFrame(df['realgdp'])
```

data

|            | realgdp   |
|------------|-----------|
| Date       |           |
| 1959-03-31 | 2710.349  |
| 1959-06-30 | 2778.801  |
| 1959-09-30 | 2775.488  |
| 1959-12-31 | 2785.204  |
| 1960-03-31 | 2847.699  |
| ...        | ...       |
| 2008-09-30 | 13324.600 |
| 2008-12-31 | 13141.920 |
| 2009-03-31 | 12925.410 |
| 2009-06-30 | 12901.504 |
| 2009-09-30 | 12990.341 |



203 rows × 1 columns

```
data.realgdp.plot(figsize=(10,6));
```




```
data['1960'].realgdp.plot(figsize=(10,6))
```

```
<ipython-input-12-3f8ea63c3e9e>:1: FutureWarning: Indexing a DataFrame with a date
data['1960'].realgdp.plot(figsize=(10,6))
<Axes: xlabel='Date'>
```



```
data['1959-03-31':'1959-12-31'].realgdp.plot(figsize=(10,6))
```

```
<Axes: xlabel='Date'>
```



```
data['1959'].mean()
```

```
<ipython-input-14-c65c702e4eab>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1959'].mean()
realgdp    2762.4605
dtype: float64
```

```
data['1960'].mean()
```

```
<ipython-input-15-69875ea1518f>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1960'].mean()
realgdp    2830.93175
dtype: float64
```

```
data['1970'].mean()
```

```
<ipython-input-16-bf22c9d7c288>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1970'].mean()
realgdp    4269.9395
dtype: float64
```

```
data['d1']=df.realgdp.diff()
```


```
data
```

|            | realgdp   | d1       |
|------------|-----------|----------|
| Date       |           |          |
| 1959-03-31 | 2710.349  | NaN      |
| 1959-06-30 | 2778.801  | 68.452   |
| 1959-09-30 | 2775.488  | -3.313   |
| 1959-12-31 | 2785.204  | 9.716    |
| 1960-03-31 | 2847.699  | 62.495   |
| ...        | ...       | ...      |
| 2008-09-30 | 13324.600 | -90.666  |
| 2008-12-31 | 13141.920 | -182.680 |
| 2009-03-31 | 12925.410 | -216.510 |
| 2009-06-30 | 12901.504 | -23.906  |
| 2009-09-30 | 12990.341 | 88.837   |

```
203 rows × 2 columns
```

```
data.d1.plot(figsize=(10,5))
```

```
<Axes: xlabel='Date'>
```



```
data['1959'].d1.mean()
```

```
<ipython-input-20-3214ad96d64c>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1959'].d1.mean()
24.95166666666667
```

```
data['1960'].d1.mean()
```

```
<ipython-input-21-c28e66afa8f8>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1960'].d1.mean()
4.3529999999999952
```

```
data['1970'].d1.mean()
```

```
<ipython-input-22-476cf48a175d>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1970'].d1.mean()
-1.6560000000001764
```

```
round(data['1959'].d1.mean())
```

```
<ipython-input-23-719979ca2937>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
round(data['1959'].d1.mean())
25
```

```
round(data['1960'].d1.mean())
```

```
<ipython-input-24-0f30aceace94>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
round(data['1960'].d1.mean())
4
```

```
round(data['1970'].d1.mean())
```

```
<ipython-input-25-a07660c039a9>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
round(data['1970'].d1.mean())
-2
```

```
data['d2']=data.realgdp.diff(periods=2)
```

data

|            | realgdp   | d1       | d2       |
|------------|-----------|----------|----------|
| Date       |           |          |          |
| 1959-03-31 | 2710.349  | NaN      | NaN      |
| 1959-06-30 | 2778.801  | 68.452   | NaN      |
| 1959-09-30 | 2775.488  | -3.313   | 65.139   |
| 1959-12-31 | 2785.204  | 9.716    | 6.403    |
| 1960-03-31 | 2847.699  | 62.495   | 72.211   |
| ...        | ...       | ...      | ...      |
| 2008-09-30 | 13324.600 | -90.666  | -42.265  |
| 2008-12-31 | 13141.920 | -182.680 | -273.346 |
| 2009-03-31 | 12925.410 | -216.510 | -399.190 |
| 2009-06-30 | 12901.504 | -23.906  | -240.416 |
| 2009-09-30 | 12990.341 | 88.837   | 64.931   |

203 rows × 3 columns

```
from pandas.core.arrays import period
data['dd1']=data.d1.diff()
```

data

|            | realgdp   | d1       | d2       | dd1      |
|------------|-----------|----------|----------|----------|
| Date       |           |          |          |          |
| 1959-03-31 | 2710.349  | NaN      | NaN      | NaN      |
| 1959-06-30 | 2778.801  | 68.452   | NaN      | NaN      |
| 1959-09-30 | 2775.488  | -3.313   | 65.139   | -71.765  |
| 1959-12-31 | 2785.204  | 9.716    | 6.403    | 13.029   |
| 1960-03-31 | 2847.699  | 62.495   | 72.211   | 52.779   |
| ...        | ...       | ...      | ...      | ...      |
| 2008-09-30 | 13324.600 | -90.666  | -42.265  | -139.067 |
| 2008-12-31 | 13141.920 | -182.680 | -273.346 | -92.014  |
| 2009-03-31 | 12925.410 | -216.510 | -399.190 | -33.830  |

```
data['1959'].dd1.mean()

<ipython-input-30-ad3b13f5a78e>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1959'].dd1.mean()
-29.36799999999971
```

```
data['1960'].dd1.mean()

<ipython-input-31-4da5448af4d0>:1: FutureWarning: Indexing a DataFrame with a datetimelike index using a single string to slice the
data['1960'].dd1.mean()
-11.530500000000075
```

```
from statsmodels.tsa.seasonal import seasonal_decompose

data.index

DatetimeIndex(['1959-03-31', '1959-06-30', '1959-09-30', '1959-12-31',
               '1960-03-31', '1960-06-30', '1960-09-30', '1960-12-31',
               '1961-03-31', '1961-06-30',
               ...,
               '2007-06-30', '2007-09-30', '2007-12-31', '2008-03-31',
               '2008-06-30', '2008-09-30', '2008-12-31', '2009-03-31',
               '2009-06-30', '2009-09-30'],
              dtype='datetime64[ns]', name='Date', length=203, freq=None)
```

```
data.realgdp.bfill(inplace=True)
```

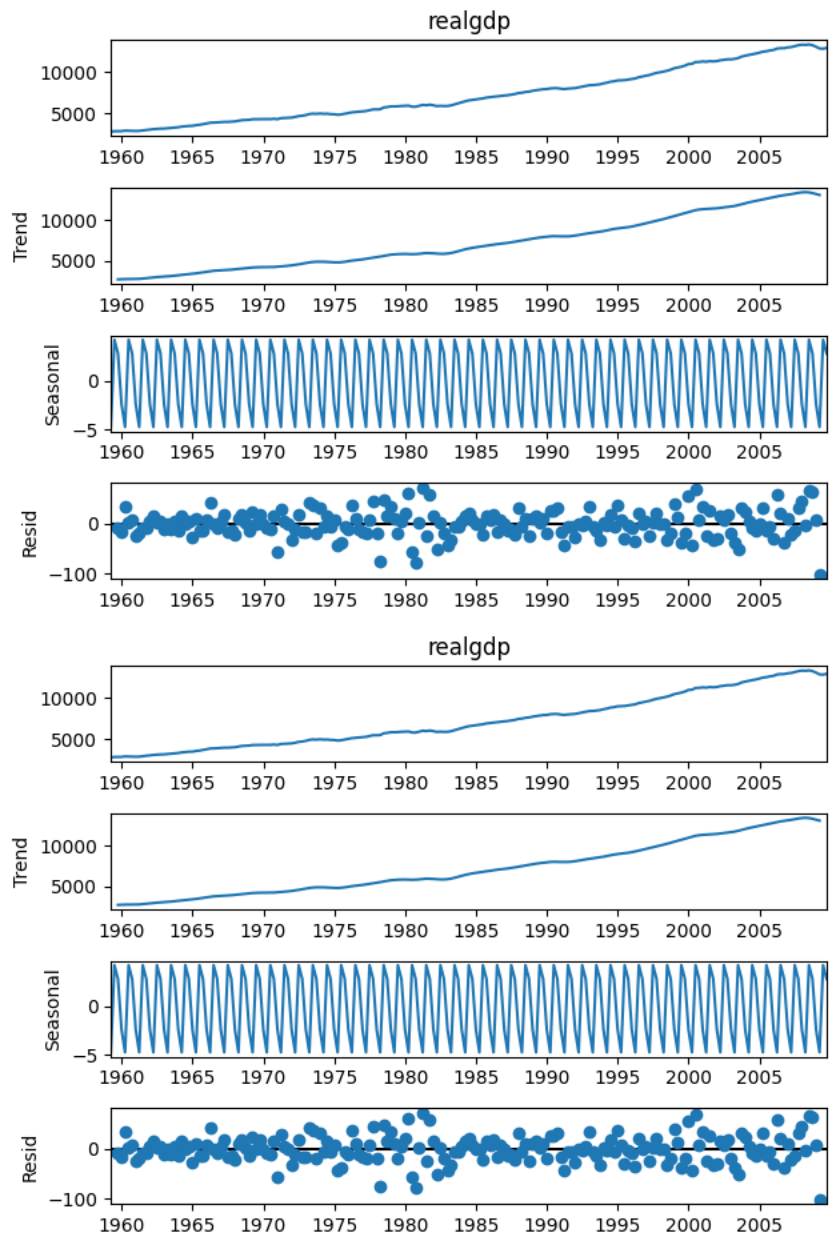
data

|            | realgdp   | d1       | d2       | dd1      |
|------------|-----------|----------|----------|----------|
| Date       |           |          |          |          |
| 1959-03-31 | 2710.349  | NaN      | NaN      | NaN      |
| 1959-06-30 | 2778.801  | 68.452   | NaN      | NaN      |
| 1959-09-30 | 2775.488  | -3.313   | 65.139   | -71.765  |
| 1959-12-31 | 2785.204  | 9.716    | 6.403    | 13.029   |
| 1960-03-31 | 2847.699  | 62.495   | 72.211   | 52.779   |
| ...        | ...       | ...      | ...      | ...      |
| 2008-09-30 | 13324.600 | -90.666  | -42.265  | -139.067 |
| 2008-12-31 | 13141.920 | -182.680 | -273.346 | -92.014  |
| 2009-03-31 | 12925.410 | -216.510 | -399.190 | -33.830  |
| 2009-06-30 | 12901.504 | -23.906  | -240.416 | 192.604  |
| 2009-09-30 | 12990.341 | 88.837   | 64.931   | 112.743  |


203 rows × 4 columns

```
x = seasonal_decompose(data.realgdp)


x.plot()
```




```
x.trend.plot()  
<Axes: xlabel='Date'>
```



```
x.seasonal.plot()
```



x.resid.plot()



data

|            | realgdp   | d1       | d2       | dd1      |
|------------|-----------|----------|----------|----------|
| Date       |           |          |          |          |
| 1959-03-31 | 2710.349  | NaN      | NaN      | NaN      |
| 1959-06-30 | 2778.801  | 68.452   | NaN      | NaN      |
| 1959-09-30 | 2775.488  | -3.313   | 65.139   | -71.765  |
| 1959-12-31 | 2785.204  | 9.716    | 6.403    | 13.029   |
| 1960-03-31 | 2847.699  | 62.495   | 72.211   | 52.779   |
| ...        | ...       | ...      | ...      | ...      |
| 2008-09-30 | 13324.600 | -90.666  | -42.265  | -139.067 |
| 2008-12-31 | 13141.920 | -182.680 | -273.346 | -92.014  |
| 2009-03-31 | 12925.410 | -216.510 | -399.190 | -33.830  |
| 2009-06-30 | 12901.504 | -23.906  | -240.416 | 192.604  |
| 2009-09-30 | 12990.341 | 88.837   | 64.931   | 112.743  |

203 rows × 4 columns

data = data.copy()

data.index



```
DatetimeIndex(['1959-03-31', '1959-06-30', '1959-09-30', '1959-12-31',
              '1960-03-31', '1960-06-30', '1960-09-30', '1960-12-31',
              '1961-03-31', '1961-06-30',
              ...
              '2007-06-30', '2007-09-30', '2007-12-31', '2008-03-31',
              '2008-06-30', '2008-09-30', '2008-12-31', '2009-03-31',
              '2009-06-30', '2009-09-30'],
              dtype='datetime64[ns]', name='Date', length=203, freq=None)
```

```
data=pd.DataFrame(data['realgdp'])
```

```
data.index
```

```
DatetimeIndex(['1959-03-31', '1959-06-30', '1959-09-30', '1959-12-31',
              '1960-03-31', '1960-06-30', '1960-09-30', '1960-12-31',
              '1961-03-31', '1961-06-30',
              ...
              '2007-06-30', '2007-09-30', '2007-12-31', '2008-03-31',
              '2008-06-30', '2008-09-30', '2008-12-31', '2009-03-31',
              '2009-06-30', '2009-09-30'],
              dtype='datetime64[ns]', name='Date', length=203, freq=None)
```

```
data = data['1959':'2009']
```

```
data.head()
```

| realgdp    |          |
|------------|----------|
| Date       |          |
| 1959-03-31 | 2710.349 |
| 1959-06-30 | 2778.801 |
| 1959-09-30 | 2775.488 |
| 1959-12-31 | 2785.204 |
| 1960-03-31 | 2847.699 |

```
Total = data.isnull().sum().sort_values(ascending = False)
Percent = (data.isnull().sum()*100/data.isnull().count()).sort_values(ascending = False)
missing_data = pd.concat([Total,Percent],axis = 1,keys=['Total','Percentage of missing value'])
missing_data
```

|         | Total | Percentage of missing value |
|---------|-------|-----------------------------|
| realgdp | 0     | 0.0                         |

```
data.plot(figsize = (15,6),legend = None)
plt.xlabel('Date',fontsize = 14)
plt.ylabel('realgdp',fontsize =14)
plt.title('Observed Monthly Average realgdp')
plt.show()
```

Observed Monthly Average realgdp



```
from statsmodels.tsa.stattools import adfuller
```

```
nnnnn |
```

```
def adf_test(timeseries):
    print('Result of Dickey_fuller Test: ')
    result=adfuller(timeseries,autolag="AIC")
    result=pd.Series(result[0:4], index=[" Test Statistic",'p-value','No. of lags Used','Number of Observations Used'])
    print(result)
```

```
if result[1]<=0.05:
    print("Strong evidence against the null hypothesis")
    print("Reject the null hypothesis")
    print("Data has no unit root and is stationary")

else:
    print("Strong evidence against the null hypothesis")
    print("Reject the null hypothesis")
    print("Data has no unit root and is stationary")
```

```
adf_test(data)
```

```
Result of Dickey_fuller Test:
  Test Statistic      1.750463
  p-value             0.998246
  No. of lags Used    12.000000
  Number of Observations Used 190.000000
  dtype: float64
  Strong evidence against the null hypothesis
  Reject the null hypothesis
  Data has no unit root and is stationary
```

```
y = data['realgdp']
```

```
y
```

```
Date
1959-03-31    2710.349
1959-06-30    2778.801
1959-09-30    2775.488
1959-12-31    2785.204
1960-03-31    2847.699
...
2008-09-30    13324.600
2008-12-31    13141.920
2009-03-31    12925.410
2009-06-30    12901.504
2009-09-30    12990.341
Name: realgdp, Length: 203, dtype: float64
```

```
train = y[:'2005']
```

```
test = y['2007':]
```

```
pip install pmdarima
```

```
Collecting pmdarima
  Downloading pmdarima-2.0.3-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux_2_28_x86_64.whl (1.8 MB)
    1.8/1.8 MB 9.1 MB/s eta 0:00:00
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.3.1)
Requirement already satisfied: Cython!=0.29.18,!=0.29.31,>=0.29 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (0.29.36)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.22.4)
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.5.3)
Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.2.2)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.10.1)
Requirement already satisfied: statsmodels>=0.13.2 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (0.13.5)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (1.26.16)
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3.10/dist-packages (from pmdarima) (67.7.2)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.19->pmdarima) (2022.7.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.22->pmdarima) (3.1.0)
Requirement already satisfied: patsy>=0.5.2 in /usr/local/lib/python3.10/dist-packages (from statsmodels>=0.13.2->pmdarima) (0.5.3)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.10/dist-packages (from statsmodels>=0.13.2->pmdarima) (23.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from patsy>=0.5.2->statsmodels>=0.13.2->pmdarima) (1.16.0)
Installing collected packages: pmdarima
Successfully installed pmdarima-2.0.3
```

```
from pmdarima.arima import auto_arima
```

```
arima_model=auto_arima(train,Seasonal=True,stepwise=False,trace=1,random_state=10)
```

```

ARIMA(0,2,0)(0,0,0)[1]      : AIC=2064.739, Time=0.07 sec
ARIMA(0,2,1)(0,0,0)[1]      : AIC=1996.294, Time=0.06 sec
ARIMA(0,2,2)(0,0,0)[1]      : AIC=1994.614, Time=0.14 sec
ARIMA(0,2,3)(0,0,0)[1]      : AIC=1986.538, Time=0.19 sec
ARIMA(0,2,4)(0,0,0)[1]      : AIC=1988.253, Time=0.31 sec
ARIMA(0,2,5)(0,0,0)[1]      : AIC=1987.307, Time=0.37 sec
ARIMA(1,2,0)(0,0,0)[1]      : AIC=2014.148, Time=0.04 sec
ARIMA(1,2,1)(0,0,0)[1]      : AIC=1990.986, Time=0.13 sec
ARIMA(1,2,2)(0,0,0)[1]      : AIC=1999.441, Time=0.22 sec
ARIMA(1,2,3)(0,0,0)[1]      : AIC=1987.740, Time=0.62 sec
ARIMA(1,2,4)(0,0,0)[1]      : AIC=1990.060, Time=0.85 sec
ARIMA(2,2,0)(0,0,0)[1]      : AIC=2011.576, Time=0.12 sec
ARIMA(2,2,1)(0,0,0)[1]      : AIC=1984.973, Time=0.43 sec
ARIMA(2,2,2)(0,0,0)[1]      : AIC=1986.046, Time=0.71 sec
ARIMA(2,2,3)(0,0,0)[1]      : AIC=inf, Time=1.54 sec
ARIMA(3,2,0)(0,0,0)[1]      : AIC=2004.129, Time=0.13 sec
ARIMA(3,2,1)(0,0,0)[1]      : AIC=1986.636, Time=0.64 sec
ARIMA(3,2,2)(0,0,0)[1]      : AIC=1987.841, Time=0.56 sec
ARIMA(4,2,0)(0,0,0)[1]      : AIC=2005.193, Time=0.08 sec
ARIMA(4,2,1)(0,0,0)[1]      : AIC=1987.898, Time=0.37 sec
ARIMA(5,2,0)(0,0,0)[1]      : AIC=2003.722, Time=0.15 sec

```


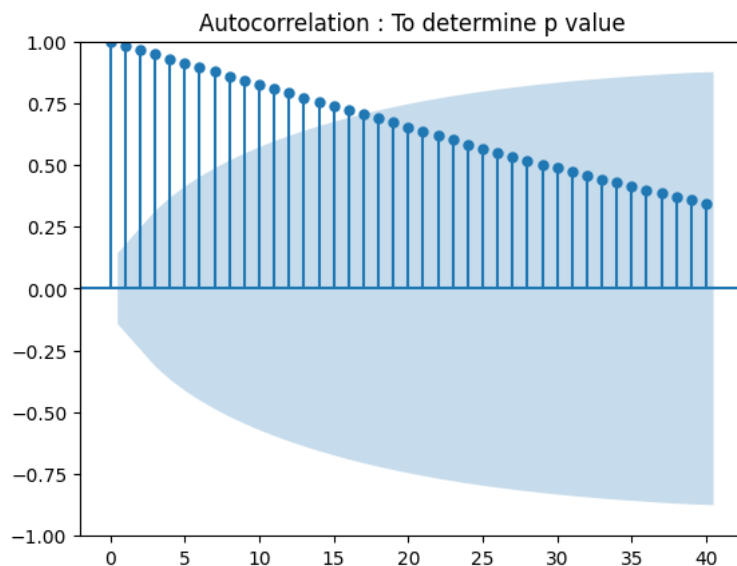
Best model: ARIMA(2,2,1)(0,0,0)[1]  
Total fit time: 7.816 seconds

```
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```

title='Autocorrelation : To determine p value'
lags=40
plot_acf(train, title=title, lags=lags)

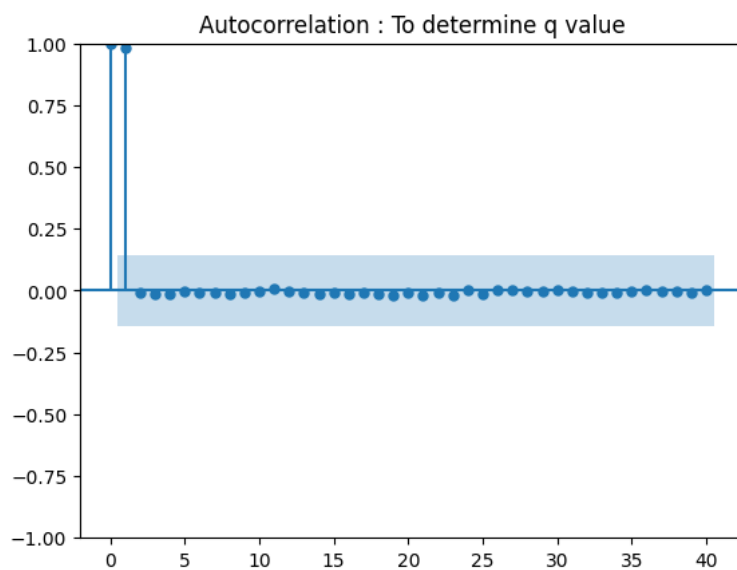
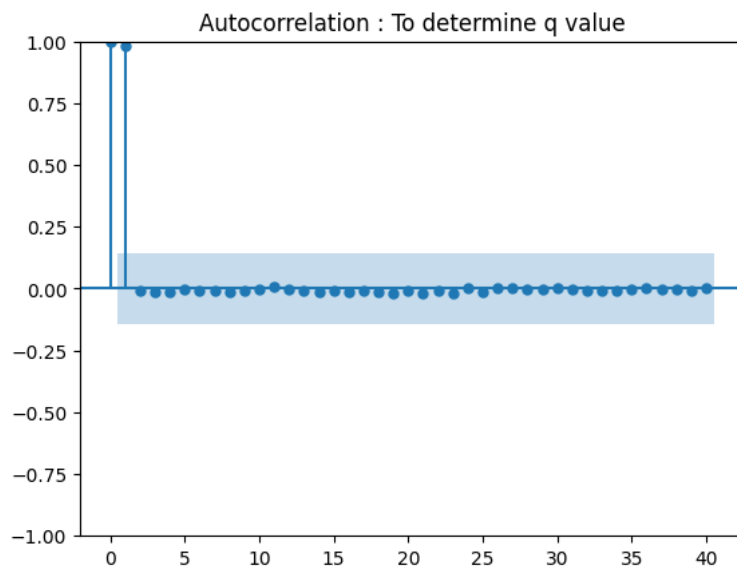
```



```

title='Autocorrelation : To determine q value'
lags=40
plot_pacf(train, title=title, lags=lags, method='ywmm')

```




```
from statsmodels.tsa.arima.model import ARIMA,ARIMAResults
```

```
model=ARIMA(train,order=(1,0,0))  
results=model.fit()  
results.summary()
```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: Val
self._init_dates(dates, freq)
```

```
model=ARIMA(train,order=(1,0,0))
results=model.fit()
results.summary()
pred=results.get_forecast(steps=36)
ax1=y['1959:'].plot(label='Observed')
pred.predicted_mean.plot(ax=ax1,label="ARIMA FORECAST",figsize=(15,6),linestyle='dashed')
pred_ci=pred.conf_int()
ax1.fill_between(pred_ci.index,pred_ci.iloc[:,0],pred_ci.iloc[:,1],color="k",alpha=0.2)
ax1.set_xlabel('Date')
ax1.set_ylabel('Average Realgdp')
plt.legend(loc= "upper left")
plt.show()
```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: Val
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: Val
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: Val
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/sarimax.py:966:
warn('Non-stationary starting autoregressive parameters')
```



```
from statsmodels.tsa.statespace.sarimax import SARIMAX
```

```
model = SARIMAX(train ,order=(1,1,2),seasonal_order=(1,0,1,12),enforce_stationarity=False,enforce_invertibility=False)
fitted_model=model.fit(maxiter=200)
print(fitted_model.summary())
```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided,
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided,
self._init_dates(dates, freq)
```

#### SARIMAX Results

```
=====
Dep. Variable:                realgdp    No. Observations:                188
Model:                SARIMAX(1, 1, 2)x(1, 0, [1], 12)    Log Likelihood                -998.489
Date:                Tue, 25 Jul 2023    AIC                2008.977
Time:                10:20:05    BIC                2028.364
Sample:                03-31-1959    HQIC                2016.833
- 12-31-2005

Covariance Type:                opg
=====
```

|          | coef      | std err  | z        | P> z  | [0.025   | 0.975]   |
|----------|-----------|----------|----------|-------|----------|----------|
| ar.L1    | 1.0000    | 0.000    | 2619.704 | 0.000 | 0.999    | 1.001    |
| ma.L1    | -0.7721   | 0.068    | -11.274  | 0.000 | -0.906   | -0.638   |
| ma.L2    | -0.1459   | 0.069    | -2.118   | 0.034 | -0.281   | -0.011   |
| ar.S.L12 | 0.8960    | 0.069    | 12.987   | 0.000 | 0.761    | 1.031    |
| ma.S.L12 | -0.9937   | 0.103    | -9.664   | 0.000 | -1.195   | -0.792   |
| sigma2   | 2434.1816 | 5.13e-05 | 4.75e+07 | 0.000 | 2434.182 | 2434.182 |

```
=====
Ljung-Box (L1) (Q):                0.09    Jarque-Bera (JB):                9.97
Prob(Q):                0.76    Prob(JB):                0.01
Heteroskedasticity (H):                1.98    Skew:                -0.12
```

Prob(H) (two-sided):0.01Kurtosis:4.10


=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 3.34e+22. Standard errors may be unstable.

```
pred=results.get_forecast(steps=36)
ax1=y[:1950:].plot(label='Observed')
pred.predicted_mean.plot(ax=ax1,label="ARIMA FORECAST",figsize=(15,6),linestyle='dashed')
pred_ci=pred.conf_int()
ax1.fill_between(pred_ci.index,pred_ci.iloc[:,0],pred_ci.iloc[:,1],color="k",alpha=0.2)
ax1.set_xlabel('Date')
ax1.set_ylabel('Average Realgdp')
plt.legend(loc= "upper left")
plt.show()
```



```
model2 = SARIMAX(train ,order=(0,0,1),seasonal_order=(1,0,1,12),enforece_stationarity=False,enforece_invertibility=False)
fitted_model=model2.fit(maxiter=200)
print(fitted_model.summary())
```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided,
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided,
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible starting MA paramete
warn('Non-invertible starting MA parameters found.')
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/sarimax.py:997: UserWarning: Non-stationary starting seasonal a
warn('Non-stationary starting seasonal autoregressive')
```

SARIMAX Results

=====

|                         |                                |                   |                   |       |          |          |
|-------------------------|--------------------------------|-------------------|-------------------|-------|----------|----------|
| Dep. Variable:          | realgdp                        | No. Observations: | 188               |       |          |          |
| Model:                  | SARIMAX(0, 0, 1)x(1, 0, 1, 12) | Log Likelihood    | -1400.619         |       |          |          |
| Date:                   | Tue, 25 Jul 2023               | AIC               | 2809.237          |       |          |          |
| Time:                   | 10:20:06                       | BIC               | 2822.183          |       |          |          |
| Sample:                 | 03-31-1959                     | HQIC              | 2814.482          |       |          |          |
|                         | - 12-31-2005                   |                   |                   |       |          |          |
| Covariance Type:        | opg                            |                   |                   |       |          |          |
| =====                   |                                |                   |                   |       |          |          |
|                         | coef                           | std err           | z                 | P> z  | [0.025   | 0.975]   |
| -----                   |                                |                   |                   |       |          |          |
| ma.L1                   | 1.0000                         | 4.089             | 0.245             | 0.807 | -7.013   | 9.013    |
| ar.S.L12                | 0.9905                         | 0.006             | 160.129           | 0.000 | 0.978    | 1.003    |
| ma.S.L12                | 0.9998                         | 4.123             | 0.242             | 0.808 | -7.081   | 9.081    |
| sigma2                  | 1.048e+05                      | 3.73e-05          | 2.81e+09          | 0.000 | 1.05e+05 | 1.05e+05 |
| =====                   |                                |                   |                   |       |          |          |
| Ljung-Box (L1) (Q):     |                                | 44.04             | Jarque-Bera (JB): |       | 2.34     |          |
| Prob(Q):                |                                | 0.00              | Prob(JB):         |       | 0.31     |          |
| Heteroskedasticity (H): |                                | 5.97              | Skew:             |       | 0.26     |          |
| Prob(H) (two-sided):    |                                | 0.00              | Kurtosis:         |       | 3.19     |          |
| =====                   |                                |                   |                   |       |          |          |

Warnings:


[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 1.79e+26. Standard errors may be unstable.

```

pred=results.get_forecast(steps=36)
ax1=y['1950:'].plot(label='Observed')
pred.predicted_mean.plot(ax=ax1,label="ARIMA FORECAST",figsize=(15,6),linestyle='dashed')
pred_ci=pred.conf_int()
ax1.fill_between(pred_ci.index,pred_ci.iloc[:,0],pred_ci.iloc[:,1],color="k",alpha=0.2)
ax1.set_xlabel('Date')
ax1.set_ylabel('Average Realgdp')
plt.legend(loc= "upper left")
plt.show()

```



```

y_forecasted_SARIMAX = pred.predicted_mean
y_truth = test
mse_SARIMAX = ((y_forecasted_SARIMAX - y_truth)**2).mean()
print('The Mean Squared Error of SARIMAX Forecast is {}'.format(round(mse_SARIMAX,2)))
print('The Root Mean squared Error of SARIMAX Forecast is {}'.format(round(np.sqrt(mse_SARIMAX),2)))

```

The Mean Squared Error of SARIMAX Forecast is 232698.8  
The Root Mean squared Error of SARIMAX Forecast is 482.39

```

model3 = SARIMAX(train ,order=(0,0,2),seasonal_order=(1,0,1,12),enforce_stationarity=False,enforce_invertibility=False)
fitted_model=model3.fit(maxiter=200)
print(fitted_model.summary())

```

```

/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided,
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:471: ValueWarning: No frequency information was provided,
self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/sarimax.py:978: UserWarning: Non-invertible starting MA parameter
warn('Non-invertible starting MA parameters found.')
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/statespace/sarimax.py:997: UserWarning: Non-stationary starting seasonal au
warn('Non-stationary starting seasonal autoregressive')

```

#### SARIMAX Results

```

=====
Dep. Variable:          realgdp      No. Observations:          188
Model:                SARIMAX(0, 0, 2)x(1, 0, [1], 12)  Log Likelihood          -1298.603
Date:                  Tue, 25 Jul 2023                AIC              2607.206
Time:                  10:20:08                        BIC              2623.388
Sample:                03-31-1959                      HQIC              2613.762
- 12-31-2005

```

```

Covariance Type:          opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ma.L1          1.0001      0.042      24.087      0.000      0.919      1.081
ma.L2          0.9989      0.077      12.916      0.000      0.847      1.150
ar.S.L12       0.9957      0.003     313.629      0.000      0.989      1.002
ma.S.L12       0.4844      0.078       6.191      0.000      0.331      0.638
sigma2       4.339e+04    3665.420      11.839      0.000    3.62e+04    5.06e+04

```

```

=====
Ljung-Box (L1) (Q):          43.36  Jarque-Bera (JB):          12.28
Prob(Q):                   0.00    Prob(JB):                   0.00
Heteroskedasticity (H):      3.94    Skew:                      0.02
Prob(H) (two-sided):         0.00    Kurtosis:                   4.25
=====

```

Warnings:  
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
y_forecasted_SARIMAX = pred.predicted_mean
y_truth = test
mse_SARIMAX = ((y_forecasted_SARIMAX - y_truth)**2).mean()
print('The Mean Squared Error of SARIMAX Forecast is {}'.format(round(mse_SARIMAX,2)))
print('The Root Mean squared Error of SARIMAX Forecast is {} '.format(round(np.sqrt(mse_SARIMAX),2)))

The Mean Squared Error of SARIMAX Forecast is 232698.8
The Root Mean squared Error of SARIMAX Forecast is 482.39
```

```
df1= pd.DataFrame(data.realgdp)
```


```
train = df1.iloc[:203]
test = df1.iloc[203:]
```

df1

| realgdp    |           |
|------------|-----------|
| Date       |           |
| 1959-03-31 | 2710.349  |
| 1959-06-30 | 2778.801  |
| 1959-09-30 | 2775.488  |
| 1959-12-31 | 2785.204  |
| 1960-03-31 | 2847.699  |
| ...        | ...       |
| 2008-09-30 | 13324.600 |
| 2008-12-31 | 13141.920 |
| 2009-03-31 | 12925.410 |
| 2009-06-30 | 12901.504 |
| 2009-09-30 | 12990.341 |

203 rows × 1 columns

```
df1.plot()
```



```
df1.size
```

203

```
df1.size-5
```

198



```
train = df1.iloc[:194]
test = df1.iloc[194:]
```

train

| realgdp    |           |
|------------|-----------|
| Date       |           |
| 1959-03-31 | 2710.349  |
| 1959-06-30 | 2778.801  |
| 1959-09-30 | 2775.488  |
| 1959-12-31 | 2785.204  |
| 1960-03-31 | 2847.699  |
| ...        | ...       |
| 2006-06-30 | 12962.462 |
| 2006-09-30 | 12965.916 |
| 2006-12-31 | 13060.679 |
| 2007-03-31 | 13099.901 |
| 2007-06-30 | 13203.977 |

194 rows × 1 columns

test.size

9

test

| realgdp    |           |
|------------|-----------|
| Date       |           |
| 2007-09-30 | 13321.109 |
| 2007-12-31 | 13391.249 |
| 2008-03-31 | 13366.865 |
| 2008-06-30 | 13415.266 |
| 2008-09-30 | 13324.600 |
| 2008-12-31 | 13141.920 |
| 2009-03-31 | 12925.410 |
| 2009-06-30 | 12901.504 |
| 2009-09-30 | 12990.341 |

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
```

```
scaler.fit(train)
```

▼ MinMaxScaler

MinMaxScaler()

```
scaled_train = scaler.transform(train)
```

```
scaled_train[:5]

array([[0.          ],
       [0.0065232  ],
       [0.00620748 ],
       [0.00713338 ],
       [0.0130889  ]])
```

```
scaled_test = scaler.transform(test)
```

```
scaled_test.max(),scaled_test.min()
```

```

(1.0201349809617797, 0.971175552988919)

from keras.preprocessing.sequence import TimeseriesGenerator

scaled_train[:5]

array([[0.          ],
       [0.0065232 ],
       [0.00620748],
       [0.00713338],
       [0.0130889 ]])

n_input = 2
n_features = 1

generator= TimeseriesGenerator(scaled_train,scaled_train,length=n_input,batch_size=1)

generator[0]

(array([[0.          ],
       [0.0065232 ]]),
 array([[0.00620748]]))

x,y = generator[0]

x

array([[0.          ],
       [0.0065232 ]])

y

array([[0.00620748]])

len(scaled_train)

194

len(generator)

192

from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
import tensorflow as tf

n_input = 5
n_features = 1
train_generator = TimeseriesGenerator(scaled_train,scaled_train,length=n_input,batch_size=1)

x.shape

(1, 2, 1)

model = tf.keras.models.Sequential([
    tf.keras.layers.LSTM(100, input_shape = (n_input,n_features), return_sequences= True),
    tf.keras.layers.LSTM(50, return_sequences = True),
    tf.keras.layers.LSTM(10),
    tf.keras.layers.Dense(64, activation = 'relu'),
    tf.keras.layers.Dense(32, activation = 'relu'),
    tf.keras.layers.Dense(1)
])

model.compile(optimizer = 'adam',loss = 'mse')

model.summary()

Model: "sequential"

```

| Layer (type) | Output Shape   | Param # |
|--------------|----------------|---------|
| lstm (LSTM)  | (None, 5, 100) | 40800   |

|                          |               |       |
|--------------------------|---------------|-------|
| lstm_1 (LSTM)            | (None, 5, 50) | 30200 |
| lstm_2 (LSTM)            | (None, 10)    | 2440  |
| dense (Dense)            | (None, 64)    | 704   |
| dense_1 (Dense)          | (None, 32)    | 2080  |
| dense_2 (Dense)          | (None, 1)     | 33    |
| =====                    |               |       |
| Total params: 76,257     |               |       |
| Trainable params: 76,257 |               |       |
| Non-trainable params: 0  |               |       |

---

CONCLUSION:

The analysis of real GDP data indicates the actual economic output. The data shows how the economy has performed over a specific period, reflecting its overall health and prosperity. By tracking real GDP, policymakers and businesses can assess the country's economic progress and make informed decisions to support future development

